

# Project 3 Overview

# Overview

- In this project you will create a filesystem
- The filesystem will support:
  - Undo logging
  - Redo logging
  - Crash recovery

# gtfs API

```
gtfs_t* gtfs_init(std::string directory, int verbose_flag);  
int gtfs_clean(gtfs_t *gtfs);  
int gtfs_clean_n_bytes(gtfs_t *gtfs, int bytes);
```

Initialization and  
on disk log  
management

```
file_t* gtfs_open_file(gtfs_t* gtfs, std::string filename, int file_length);  
int gtfs_close_file(gtfs_t* gtfs, file_t* fl);  
int gtfs_remove_file(gtfs_t* gtfs, file_t* fl);
```

File  
creation/removal

```
char* gtfs_read_file(gtfs_t* gtfs, file_t* fl, int offset, int length);  
write_t* gtfs_write_file(gtfs_t* gtfs, file_t* fl, int offset, int length, const char* data);  
int gtfs_sync_write_file(write_t* write_id);  
int gtfs_sync_write_file_n_bytes(write_t* write_id, int bytes);  
int gtfs_abort_write_file(write_t* write_id);
```

Reading and  
Writing data

Write Visibility Scenarios

# A process can read its own writes, without syncing

```
int main(){  
    gtfs_init(dir_abs_path)  
  
    gtfs_open_file("file1.txt", length)  
  
    gtfs_write_file("hello1", offset, n_bytes)  
  
    data * = gtfs_read_file(offset, n_bytes)  
  
    assert(strcmp(data, "hello1") == 0)  
  
    return 0  
  
}
```

# A process can abort a write, and read a version of the file as if that write didn't happen

```
int main(){  
  
    gtfs_init(dir_abs_path)  
  
    gtfs_open_file("file1.txt", length)  
  
    wr1 = gtfs_write_file("hello1", offset, n_bytes)  
  
    wr2 = gtfs_write_file("hello2", offset, n_bytes) //overwrite  
  
    data * = gtfs_read_file(offset, n_bytes)  
  
    assert(strcmp(data, "hello2") == 0)  
  
    gtfs_abort_write(wr2)  
  
    data * = gtfs_read_file(offset, n_bytes)  
  
    assert(strcmp(data, "hello1") == 0) //read the first write at this offset  
  
    return 0  
  
}
```

A process's write only becomes visible to other processes, if the write is sync'd.

```
//process 1
```

```
int main(){
```

```
    gtfs_init(dir_abs_path)
```

```
    gtfs_open_file("file1.txt", length)
```

```
    wr1= gtfs_write_file("hello1", offset, n_bytes)
```

```
    gtfs_sync_write_file(wr1) //added to on disk log
```

```
    return 0
```

```
}
```

A process's write only becomes visible to other processes, if the write is sync'd.

//process 1

```
int main(){  
    gtfs_init(dir_abs_path)  
    gtfs_open_file("file1.txt", length)  
    wr1= gtfs_write_file("hello1", offset, n_bytes)  
    gtfs_sync_write_file(wr1) //added to on disk log  
  
    return 0  
}
```

//process 1 has exited, process 2 now

```
int main(){  
    gtfs_init(dir_abs_path)  
    gtfs_open_file("file1.txt", length)  
    data * = gtfs_read_file(offset, n_bytes)  
    assert(strcmp(data, "hello1") == 0)  
    return 0  
}
```



# A process's write that isn't sync'd will not be visible from a second process

```
//process 1
```

```
int main(){
```

```
    gtfs_init(dir_abs_path)
```

```
    gtfs_open_file("file1.txt", length)
```

```
    wr1= gtfs_write_file("hello1", offset, n_bytes)
```

```
    //we did not call gtfs_sync_write
```

```
    return 0
```

```
}
```

# A process's write that isn't sync'd will not be visible from a second process

//process 1

```
int main(){
```

```
    gtfs_init(dir_abs_path)
```

```
    gtfs_open_file("file1.txt", length)
```

```
    wr1= gtfs_write_file("hello1", offset, n_bytes)
```

```
    //we did not call gtfs_sync_write
```

```
    return 0
```

```
}
```

//process 1 has exited, process 2 now

```
int main(){
```

```
    gtfs_init(dir_abs_path)
```

```
    gtfs_open_file("file1.txt", length)
```

```
    data * = gtfs_read_file(offset, n_bytes)
```

```
    assert(strcmp(data, "") == 0) //uninitialized reads return "". See API spec
```

```
    return 0
```

```
}
```

Log management on disk / in memory

Doing a clean(), should apply the on disk log to the file,  
and truncate the log size to 0

```
total 4
-rwxrwxr-x 1 daniel daniel 100 Mar  9 10:38 additional1.txt
-rwxrwxr-x 1 daniel daniel  28 Mar  9 10:38 additional1.txt.log
total 4
-rwxrwxr-x 1 daniel daniel 100 Mar  9 10:38 additional1.txt
-rwxrwxr-x 1 daniel daniel   0 Mar  9 10:38 additional1.txt.log
```

gtfs\_clean() //applies and truncates log  
Reflected in file size

# Test Cases

- The starter code comes with test cases to check for expected filesystem behavior
  - Write to file, then read that data back from the same process
  - Write to file, then overwrite that write, abort the latter write, then read the first write from the same process
  - Write to file from one process:
    - Then sync that write, and read back the same data from a second process
    - Then abort that write, and check that the data cannot be read from a second process
  - Checks for on-disk log appending: syncing writes should increase the on disk size of redo-log
  - Checks for log compaction: cleaning the on-disk log should reduce their size (either to 0, or by some partial amount if using the `clean_n_bytes()` api)

# Other Notes

- You can assume that only one process at a time will have the file system open
  - That is, when a process calls 'gtfs\_init' you can assume as a precondition that no other running process has issued this API call
- Uninitialized reads that are in file bounds return ""
  - Not the best interface design, but the tests assume this